

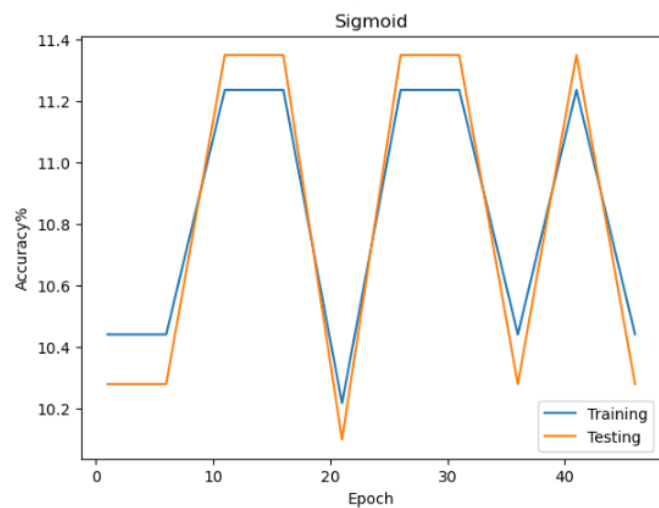

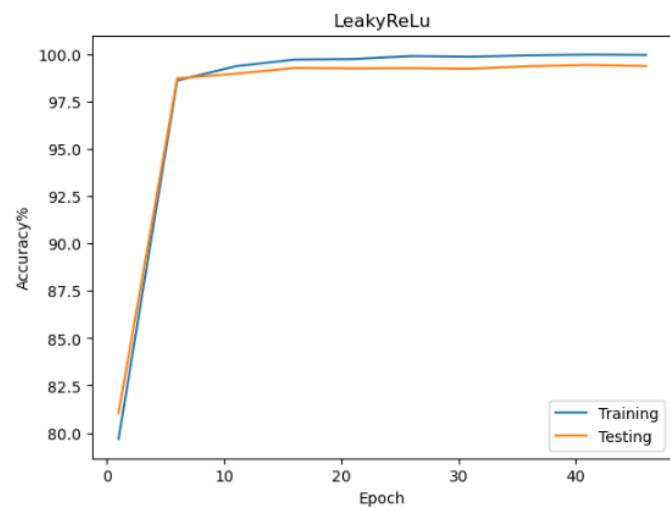
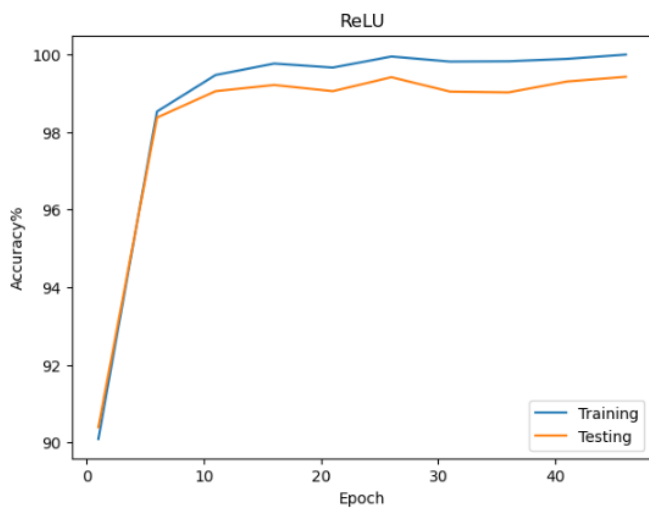
Assignment: 2

Q1

a) (Which Activation function is better?)

- only RELU activation
- only leaky RELU activation,
- only sigmoid activation,
- only tanh activation.

ReLU performed the best out of all the activation functions, LeakyReLU and Tanh performed similarly only sigmoid wasn't able to learn.



Optimizer: Adam learning rate 0.0001

Loss Function: Cross Entropy Loss





b) (What learning rate is better?)

Optimizer: Adam

Loss Function: Cross Entropy Loss

Activation Function: ReLU

Best learning rate was found to be 0.00001



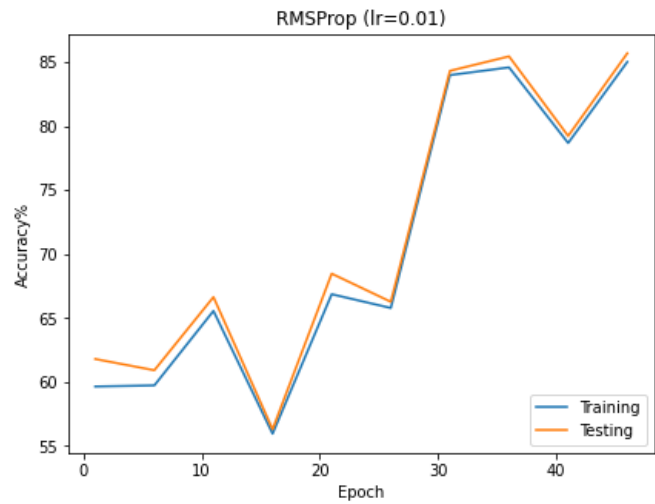
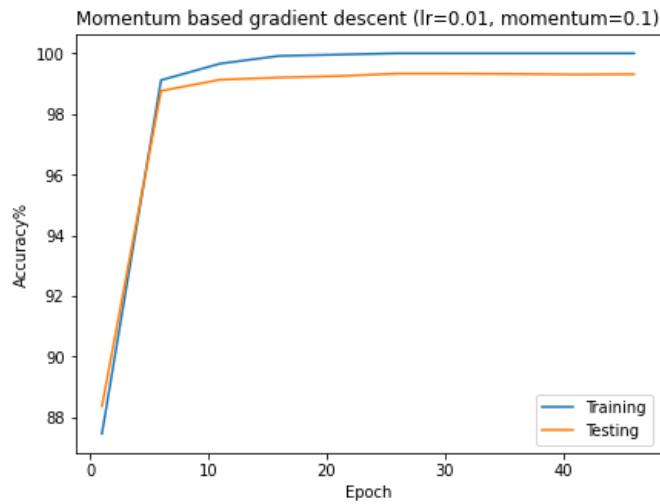
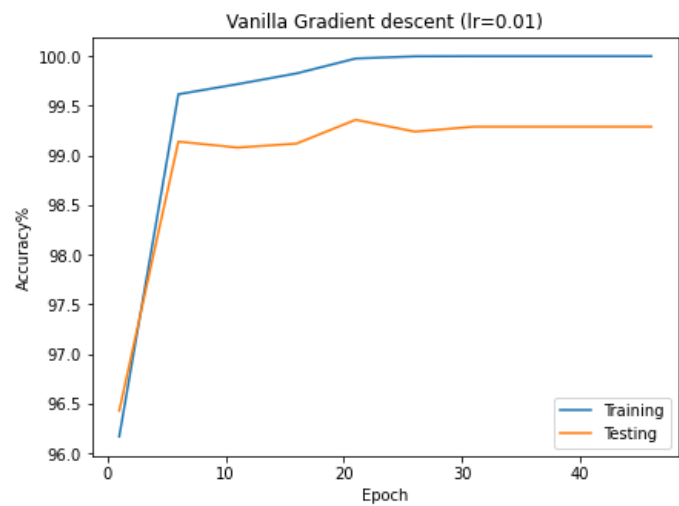

c) (Role of optimizer)

- Vanilla Gradient descent
- Momentum based gradient descent (momentum value of your choice)
- Adam.
- RMS prop.

Loss Function: Cross Entropy Loss

Activation Function: ReLU

Gradient descent both Vanilla and momentum based performed comparatively better.



d) (Transfer Learning)


Optimal hyperparameters:

Optimizer: Momentum based gradient descent (lr: 0.01, momentum: 0.1)

Loss Function: Cross Entropy Loss

Activation Function: Tanh

Learning on the pretrained VGG16 model {optimizer: SGD with momentum, lr: 0.1, loss: crossEntropy, activation: Tanh}



Q2




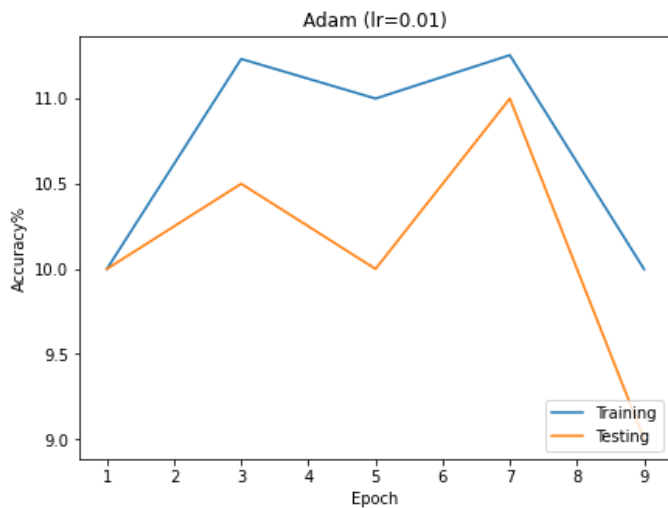
- a) Initialize the ResNet-18 network with pre-trained weights from ImageNet and then try to use these weights to improve the training for the CIFAR-10 dataset.

hyperparameters:

Optimizer: Adam

Loss Function: Cross Entropy Loss

Activation Function: Tanh



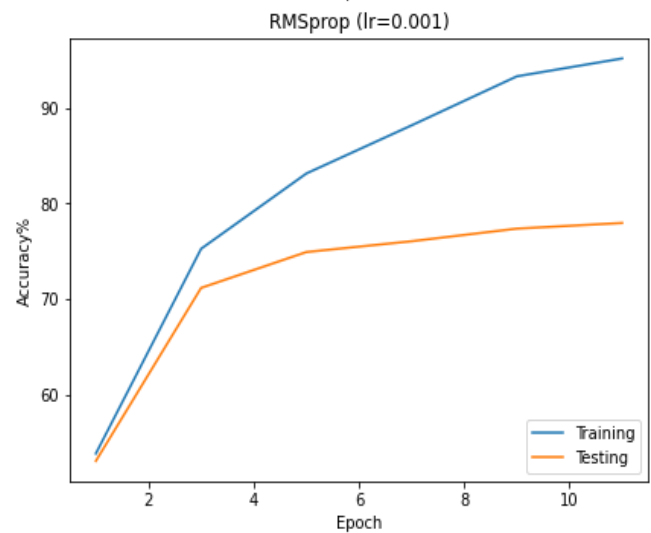
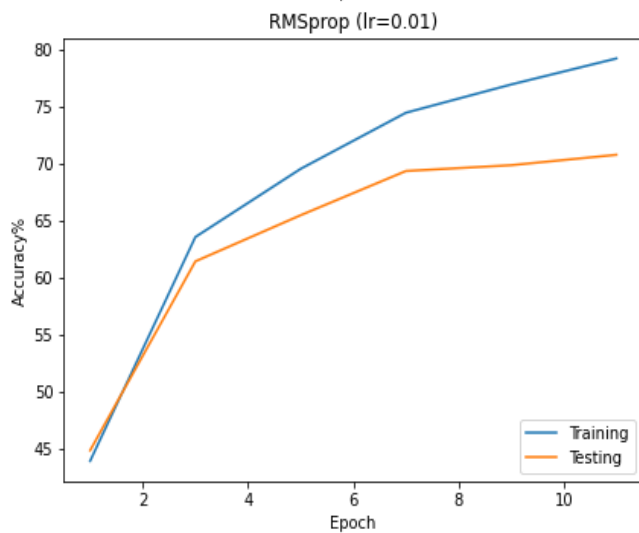

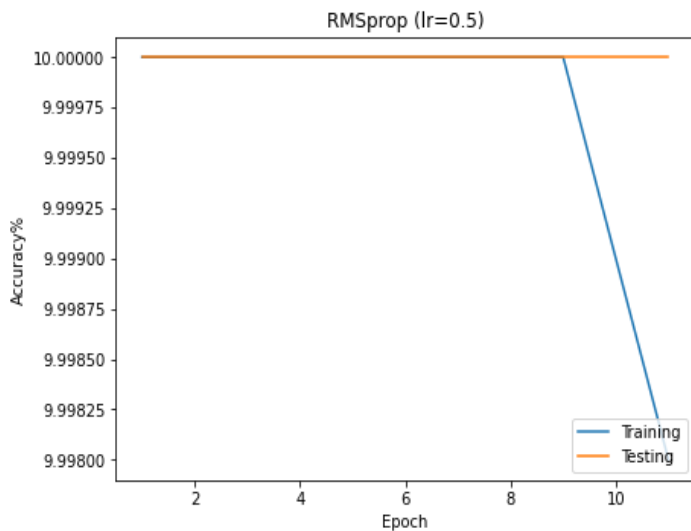
Best performing in Adam(lr=0.0001) Accuracy: 78.26%

hyperparameters:

Optimizer: RMSprop

Loss Function: Cross Entropy Loss

Activation Function: ReLU



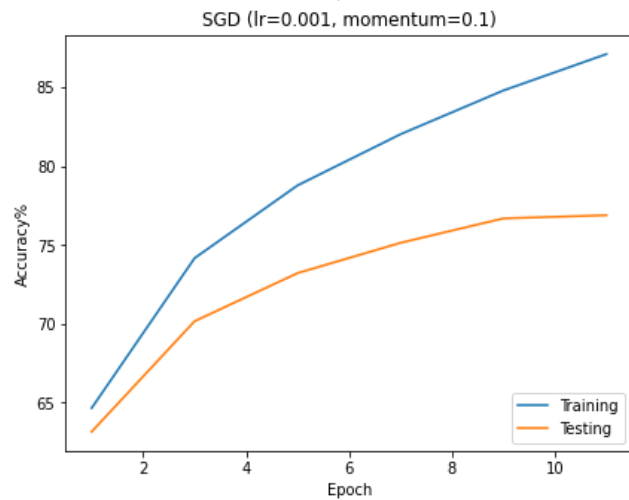



Best performing in RMSprop(lr=0.001) Accuracy: 77.92%

hyperparameters:

Optimizer: Momentum based gradient descent (momentum: 0.1)

Loss Function: Cross Entropy Loss

Activation Function: ReLU



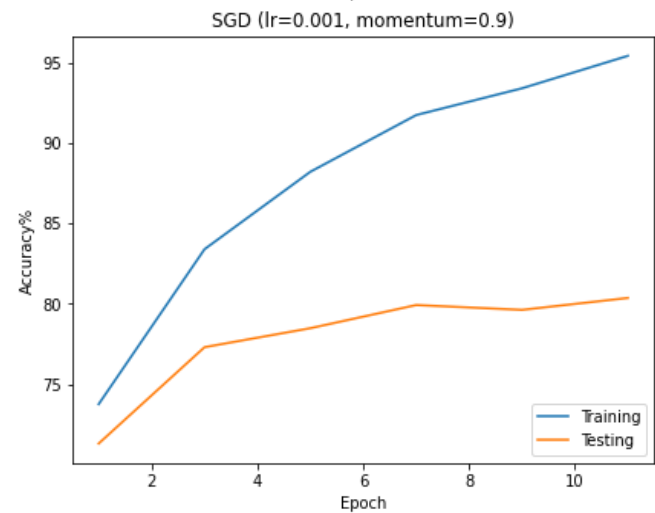



Best performing in SGD(lr=0.01, momentum=0.1) Accuracy: 80.15%

hyperparameters:

Optimizer: Momentum based gradient descent (momentum: 0.9)

Loss Function: Cross Entropy Loss

Activation Function: ReLU



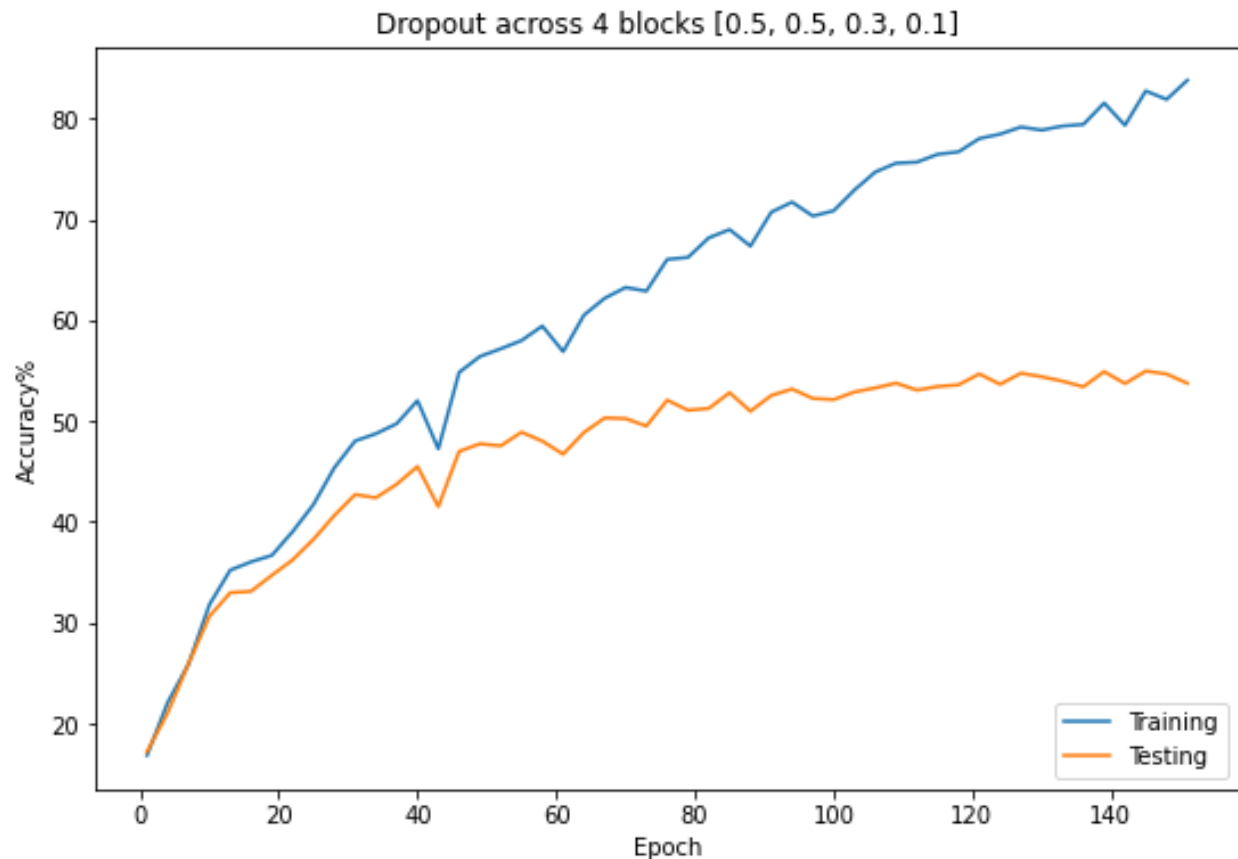
Best performing in SGD(lr=0.01, momentum=0.9) Accuracy: 79.88%

- b) Train the network from scratch with the Tiny-CIFAR-10 dataset. Try different data augmentations and dropouts after different layers and with different dropout rates.


Following augmentations are used

```
transforms.RandomAffine(degrees=(30, 70)),  
transforms.RandomPerspective(distortion_scale=0.6, p=1.0),  
transforms.ElasticTransform(alpha=150.0),  
transforms.RandomSolarize(threshold=192.0),  
transforms.ToTensor(),  
transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
```


Result: however different the dropout rates network converges with test accuracy around 53-55%





Dropout across 4 blocks [0.5, 0.4, 0.3, 0.2]





Dropout across 4 blocks [0.1, 0.2, 0.3, 0.4]







c) Visualize the activations of the CNN for a few test examples in each of the above cases.



1) Dropout across 4 blocks [0.5, 0.5, 0.3, 0.1]



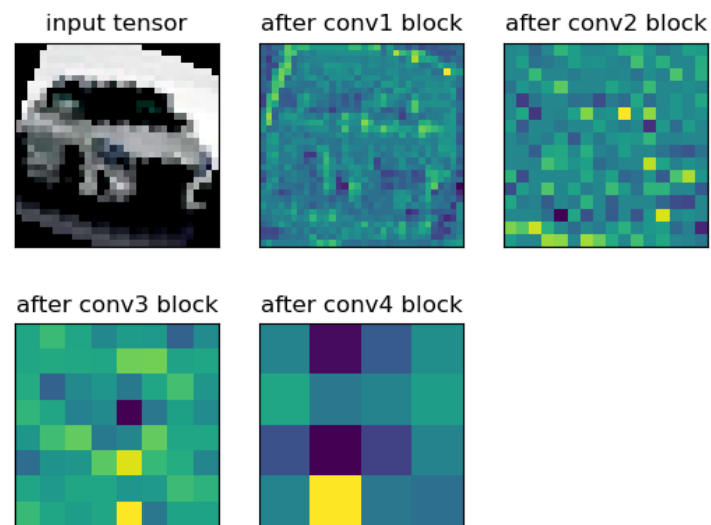

2) Dropout across 4 blocks [0.5, 0.4, 0.3, 0.2]



3) Dropout across 4 blocks [0.1, 0.2, 0.3, 0.4]



4) Dropout across 4 blocks [0.3, 0.3, 0.3, 0.3]



Results: It is visible that in the first few layers focus of the network is more towards borders whereas in the later layers it focuses on more complex features which are in the most part unintelligible for human perception.